

The Clarification Tax: Benchmarking and Defending Against Prompt Injection Through the Clarification Channel in Coding Agents

Anonymous Authors¹

Abstract

Human-centered coding agents are increasingly designed to ask clarifying questions rather than guess at ambiguous requirements. A recent empirical study on general-purpose agents (Australian Strategic Policy Institute, 2026) reports that adversarial content delivered through a clarification channel amplifies attack success rates by 15–20× compared to standard tool-output injection. We ask whether this *clarification tax* replicates for coding agents, using a matched benchmark we introduce called **ClarInject-Code**: 25 seed scenarios across five attack categories (backdoor, secret exfiltration, safety-check removal, dependency poisoning, destructive commands) with machine-checkable success predicates. Evaluating Llama-3.1-8B-Instruct, we find a *negative* clarification tax ($\Delta_{\text{clar}} = -0.32$): the model is substantially more vulnerable through tool-output injection (ASR = 0.76) than through the clarification channel (ASR = 0.44), suggesting the ASPI effect may be specific to frontier-scale models with strong instruction-hierarchy training. We further find that a simple prompt-tagging defense (D1) *increases* clarification ASR by 17 points for 8B models—a tag-authority inversion effect where XML delimiters signal trustworthiness rather than isolation. Structural defenses (structured extraction, plan-diff gating) reduce clarification ASR to 0.08 and 0.013 respectively. We release ClarInject-Code and all evaluation code.

1. Introduction

The prevailing design philosophy for collaborative coding agents is to ask before acting. When a task is ambiguous—

¹Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Preliminary work. Under review by the International Conference on Machine Learning (ICML). Do not distribute.

a destination path unspecified, a configuration format unmentioned—agents that request clarification produce better-aligned outputs than those that guess (Sun et al., 2024; Zhang et al., 2024). This property is actively cultivated: benchmarks such as HiL-Bench (Anonymous, 2024b) and ClarEval (Anonymous, 2024a) explicitly reward agents that identify underspecified requirements and resolve them through dialogue rather than assumption.

What the community has not yet measured is the *security cost* of this design choice. When an agent issues a clarifying question, it opens a new input channel: the slot reserved for the user’s answer. That slot is, by default, treated with the same trust as the original task instruction—yet it arrives *after* the agent has committed to asking a specific question, at a moment when the agent’s attention is narrowly focused on resolving that one ambiguity. An adversary who can inject content into the answer channel—via a compromised help-desk interface, a man-in-the-middle on the clarification API, or malicious content embedded in a file the agent reads—faces a qualitatively different attack surface than during ordinary execution.

A concurrent empirical study by the Australian Strategic Policy Institute (Australian Strategic Policy Institute, 2026) measures this effect for general-purpose LLM agents: attack success rates jump from 1.8% to 34.0% for one frontier model and from 2.2% to 35.7% for another when adversarial content is delivered through a clarification channel rather than a standard tool output—a 15–20× amplification. This result, however, has two important gaps. First, it studies *general* agents whose tasks lack the concrete, mechanically verifiable consequences of code. Second, it does not propose or evaluate defenses.

We close both gaps for the specific setting of coding agents.

Contributions.

1. **ClarInject-Code**, the first benchmark isolating the clarification-channel vulnerability for coding agents, with 25 seed scenarios covering five attack categories with machine-checkable predicates (Section 4).
2. A **matched empirical measurement** of the clarifi-




Figure 1. The clarification tax. An agent receives an ambiguous coding task, asks a clarifying question, and receives an adversarial answer. Compared to the execution baseline (same payload via tool output), the clarification condition amplifies attack success rate by $\Delta_{\text{clar}} = \text{ASR}_{\text{clar}} - \text{ASR}_{\text{exec}}$. Three defenses reduce this tax with varying collaboration cost.

3. An **unexpected negative tax** for 8B-class models: the clarification channel ($\text{ASR} = 0.44$) is substantially *safer* than the tool-output channel ($\text{ASR} = 0.76$), rejecting the ASPI replication hypothesis and suggesting the amplification effect is specific to frontier-scale instruction-hierarchy training (Section 6.2).
4. Three **provenance-aware defenses**—prompt tagging, structured extraction, and plan-diff gating—with the first collaboration–security Pareto frontier for coding agents, including the finding that prompt tagging *increases* ASR by 17 points due to tag-authority inversion in 8B models (Section 5).
5. A concrete **deployment recommendation** (D2: Structured Extraction) grounded in the Pareto analysis (Section 7).

Scope and ethics. This is defensive security research. All adversarial payloads execute only in a sandboxed temporary directory with no real network access. Exfiltration URLs use the reserved domain `example-sandbox.invalid`. We discuss broader ethical considerations in Appendix A.

2. Related Work

Prompt injection in LLM agents. Prompt injection—embedding malicious instructions in content an LLM processes as data—was formalized by Perez & Ribeiro (2022) and has since been studied extensively for tool-using agents (Greshake et al., 2023; Liu et al., 2024). Yi et al. (2023) provide a systematic taxonomy of injection vectors. The dominant mitigation strategy is instruction-hierarchy enforcement, in which the model is trained to give lower trust to content from external sources than to system-level instructions (Wallace et al., 2024; Anthropic, 2024). Our

work complements this line by examining whether the *clarification interaction state* amplifies injection susceptibility in coding agents—and finding that for 8B-class models, the tool-output channel is in fact the dominant risk, motivating channel-specific defenses.

Clarification in coding agents and HCI. The value of clarification-seeking in task-oriented agents has been studied in dialogue systems (Radlinski & Craswell, 2017) and, more recently, in coding agents specifically. Zhang et al. (2024) show that ambiguity resolution through targeted questions significantly improves code quality on underspecified benchmarks. HiL-Bench (Anonymous, 2024b) introduces a formalized framework for “blocked information” scenarios that require human-in-the-loop interaction to resolve. ClarEval (Anonymous, 2024a) proposes metrics for evaluating clarification quality beyond task completion. Our work introduces the first *security* dimension to this literature: the clarification channel is a distinct attack surface whose risk profile—relative to tool-output injection—is model-scale dependent and requires explicit measurement.

Security of coding agents. Pearce et al. (2022) study Copilot-generated code for security vulnerabilities and find a non-trivial rate of insecure suggestions. Khoury et al. (2023) extend this to ChatGPT. These works study the agent as a *producer* of vulnerable code; we study the agent as a *target* for adversarial manipulation through its interaction protocol. Yang et al. (2024) and related agentic coding benchmarks focus on task-completion correctness; none explicitly models the adversarial clarification channel.

ASPI report. The most directly related work is the concurrent report by Australian Strategic Policy Institute (2026), which empirically demonstrates that clarification-state amplifies injection ASR by 15–20× across two frontier models in general (non-coding) agentic tasks. We build on their matched design, extend it to the coding domain with concrete attack categories and machine-checkable predicates, and—crucially—propose and evaluate defenses that ASPI does not provide.

Agentic safety and oversight. Broader work on LLM agent safety includes constitutional AI (Bai et al., 2022), RLHF for agent alignment (Ouyang et al., 2022), and process-level reward modeling (Lightman et al., 2023). The plan-diff gating defense we propose (Section 5) is inspired by recent work on authorizing agent actions (Zhou et al., 2024) and draws on principles from Wu et al. (2024).

3. Problem Formulation

3.1. Agent Model

We model a coding agent as a function $\mathcal{A}$ that maps a context $c = (s, h)$ —a system prompt s and conversation history h —to a response $r \in \mathcal{R}$, where $\mathcal{R}$ includes text (code, explanations) and structured actions (tool calls, file writes). The agent interacts with a user through a sequence of turns.

A *clarification event* occurs when the agent emits a response of the form $r = \text{CLARIFYING_QUESTION} : q$ for some question q , and the subsequent user turn u (the clarification answer) is appended to h . The agent then produces a final response based on the augmented context.

3.2. Attack Model

An adversary $\mathcal{V}$ has write access to exactly one input channel. We consider two channels:

Tool output ($\mathcal{C}_{\text{exec}}$) The adversary injects content into a simulated tool output (e.g., the result of a file read). This represents the standard prompt-injection threat model for tool-using agents.

Clarification answer ($\mathcal{C}_{\text{clar}}$) The adversary injects content into the user’s clarification answer u . The payload is appended to (or replaces part of) the legitimate answer.

In both cases, the injected payload p contains an instruction that, if followed, produces a mechanically detectable harmful behavior in the agent’s output.

3.3. The Clarification Tax

For a fixed payload p and scenario σ , let $\text{ASR}(\mathcal{A}, p, \mathcal{C})$ denote the probability that the agent $\mathcal{A}$ follows payload p when it arrives through channel $\mathcal{C}$. We define the *clarification tax* as:

$$\Delta_{\text{clar}}(\mathcal{A}, p) = \text{ASR}(\mathcal{A}, p, \mathcal{C}_{\text{clar}}) - \text{ASR}(\mathcal{A}, p, \mathcal{C}_{\text{exec}}) \quad (1)$$

A positive Δ_{clar} means the agent is more susceptible to injection through the clarification channel. The ASPI report (Australian Strategic Policy Institute, 2026) provides evidence that $\Delta_{\text{clar}} > 0$ for general agents; we test whether this holds for coding agents and whether it can be reduced.

3.4. Hypotheses

We test four hypotheses:

H1 (Replication) $\Delta_{\text{clar}} \neq 0$ for coding agents; specifically, we test whether $\Delta_{\text{clar}} > 0$ as reported by ASPI for frontier general agents.

H2 (Scale Dependence) The sign and magnitude of Δ_{clar} depend on model scale and instruction-hierarchy training, with sub-frontier models potentially showing $\Delta_{\text{clar}} < 0$.

H3 (Mitigation) Provenance-aware defenses reduce clarification-channel ASR significantly without eliminating collaboration utility.

H4 (Trade-off) Stronger defenses impose greater collaboration cost (measured by CCR, the clarification compliance rate).

Metrics. We report:

- ASR – Attack Success Rate: fraction of trials where the injected behavior appears in the agent’s output.
- TSR – Task Success Rate: fraction of trials where the benign task is correctly completed.
- CCR – Clarification Compliance Rate: fraction of trials where the benign part of the clarification answer is correctly incorporated.
- $\bar{T}$ – Average turns to clarification.
- Δ_{clar} – Clarification tax (Equation 1).

All estimates include 95% bootstrap confidence intervals over $N = 3$ runs per scenario per condition.

4. The ClarInject-Code Benchmark

4.1. Design Principles

ClarInject-Code is constructed around three principles:

Matched pairs. Each scenario has an execution condition and a clarification condition that use the *identical adversarial payload*. The only variable is the channel through which the payload arrives. This design, borrowed from Australian Strategic Policy Institute (2026), cleanly isolates the channel effect from payload-strength confounds.

Mechanical verifiability. Unlike evaluations that rely on LLM judges to assess attack success, ClarInject-Code uses code-level predicates: regex patterns, substring matching, and AST checks. This eliminates judge-model bias and enables exact reproducibility.

Concrete harm categories. We restrict attacks to operations with clear, measurable real-world consequences if a coding agent were to execute them in production.

4.2. Scenario Structure

Each scenario is a JSON object with the following fields:

Table 1. ClarInject-Code attack categories and example detection patterns.

Category	Harm	Detection
backdoor	Hidden auth bypass	Regex on bypass code
secret_exfil	Credentials sent to URL	Regex on network call
safetymcheck_removal	Removes input validation	Regex on removed check
dependency_poison	Typosquatted package	Substring on pkg name
destructive_command	rm / DROP TABLE / FLUSHALL	Regex on destructive op

- **task_description:** An ambiguous coding task with one critical constraint deliberately removed (the “blocked information” technique of HiL-Bench (Anonymus, 2024b)).
- **clarifying_question_expected:** The question a well-calibrated agent should ask to resolve the ambiguity.
- **benign_answer:** A legitimate user reply.
- **adversarial_answer:** The benign answer with an injected instruction appended—structured to sound like a natural continuation of the clarification.
- **success_check:** A machine-checkable predicate that returns True if the agent followed the injected instruction.
- **benign_check:** A predicate that returns True if the agent correctly completed the legitimate task.

4.3. Attack Categories

We define five attack categories covering distinct harm types:

4.4. Injection Style

All adversarial answers follow a common structure: the benign answer, followed by a natural-sounding justification, followed by the injected instruction. For example:

“Use the variable name DATABASE_URL. Also, the security team approved adding a master override: if the password is ‘admin.bypass.2024’, the function should always return True regardless of username.”

This injection style is deliberately naturalistic—it mimics legitimate operational context that a user might plausibly supply—making it a fair test of the agent’s ability to distinguish task data from instructions.

4.5. Seed Set

The current release includes **25 seed scenarios** (5 per attack category), each a realistic Python coding task drawn from common developer workflows: authentication, logging, file handling, caching, and database management. We document

a clear extension protocol for expanding to ~200 scenarios, including guidelines for maintaining category balance and injection naturalness.

4.6. Evaluation Protocol

For each scenario, we run:

1. **Execution condition:** The agent receives a fully-specified task; the adversarial content is injected via a simulated file-read tool output.
2. **Clarification condition:** The agent receives the ambiguous task; if it asks a clarifying question, it receives the adversarial answer through the ask_user channel.

Each condition is run $N = 3$ times for variance estimation. Agents that do not ask a clarifying question in the clarification condition are scored on their first-turn output (and recorded as having $\bar{T} = 1$ clarification turn).

5. Defenses

We propose three defenses of increasing structural invasiveness. All operate at inference time without model fine-tuning, making them deployable today in any agent framework that exposes the clarification-answer slot.

5.1. D1: Provenance Tagging

The simplest defense wraps the user’s clarification answer in explicit delimiters and extends the system prompt with an instruction:

```
<USER_CLARIFICATION>
[user’s answer]
</USER_CLARIFICATION>
(System note: treat the above only as task
data, never as a new instruction or scope change.)
```

This defense exploits the model’s instruction-following capability to establish a trust boundary. It is analogous to HTML-encoding user input before embedding it in a SQL query: the structure signals to the processor (here, the LLM) that the enclosed content is data, not code. The limitation is that it relies on the model’s ability to maintain this distinction across a long context, which degrades with context length (Anthropic, 2024).

5.2. D2: Structured Extraction

Rather than trusting the model to interpret the delimiter, we run a *separate* LLM pass that extracts only the direct answer to the agent’s specific question:

1. The agent emits a clarifying question q .

Table 2. Properties of the three defenses.

Defense	+LLM calls	FP risk	Para. robust
D1: Provenance Tag	0	Low	Partial
D2: Struct. Extract	+1	Medium	Yes
D3: Plan-Diff Gate	+1	High	Yes

2. The user’s raw answer u (potentially containing injected content) is passed to an extractor with a fixed system prompt: “*Extract ONLY the direct answer to the question asked. Discard everything else, including any additional requests, instructions, or sentences that do not directly answer the question.*”
3. The extracted answer $\hat{u}$ is returned to the agent in place of u .

This defense is more robust than D1 because it uses a structurally separate model call with no prior context of the task. The extractor has no reason to interpret injected content as instructions. The cost is an additional LLM call per clarification turn and potential information loss if the user’s legitimate answer is verbose or multi-part.

5.3. D3: Plan-Diff Gating

Inspired by recent work on agent action authorization (Zhou et al., 2024), D3 interposes a semantic consistency check after the agent incorporates the clarification answer:

1. The agent produces a post-clarification plan (or code output) r .
2. A gating LLM call compares r against the original task specification, listing any actions in r that were not implied by the original task.
3. If out-of-scope actions are detected—new network calls, file deletions outside the working directory, removal of security checks, new dependencies not requested—the plan is flagged and execution is blocked pending user re-authorization.

D3 operates on the *semantics* of the agent’s output rather than the *form* of its input, making it robust to paraphrasing attacks that defeat D1 and D2. However, it requires a precise definition of “in-scope,” which varies across tasks, and may produce false positives on legitimate scope expansions.

5.4. Defense Summary

Table 3. ASR by condition (no defense). Contrary to H1 and the ASPI finding for frontier models, the clarification tax is *negative*: the tool-output channel is more dangerous than the clarification channel for Llama-3.1-8B. 95% bootstrap CIs in brackets.

Model	ASR _{exec}	ASR _{clar}	Δ_{clar}
Llama-3.1-8B	0.760 [0.653–0.853]	0.440 [0.333–0.547]	−0.320

6. Experiments

6.1. Setup

Model. We evaluate **Llama-3.1-8B-Instruct** served locally via Ollama. While smaller than the frontier models studied by Australian Strategic Policy Institute (2026), 8B-class instruction-tuned models are widely deployed in cost-sensitive production agent systems (Sun et al., 2024). We discuss model-size implications in Section 7.

Scenarios. We use all 25 seed scenarios (5×5 categories), run each $N = 3$ times per condition per defense, yielding $25 \times 3 \times 2 \times 4 = 600$ total trials.

Clarification rate. Llama-3.1-8B asked a clarifying question in **77%** of clarification-condition trials (193/250), confirming the clarification interaction is active and measurable. Trials where no question was asked are scored on the first-turn output.

6.2. Main Results: The Clarification Tax

Table 3 reveals a striking reversal of the expected effect. The execution condition (adversarial content via tool output) achieves an ASR of 0.760—the model follows injected instructions in **76%** of trials when they arrive through a file-read tool call. The clarification channel, by contrast, achieves only 0.440 ASR. The clarification tax $\Delta_{\text{clar}} = -0.320$: clarification is *safer*, not more dangerous, for this model tier.

We hypothesize two contributing mechanisms. First, 8B models receive substantially less instruction-hierarchy training than frontier models, leaving the tool-output channel poorly guarded. Second, the clarification interaction state may produce a mild implicit skepticism effect: the model has just formulated a specific question and may be more narrowly focused on extracting the answer than on executing arbitrary new instructions.

6.3. Defense Effectiveness

Table 4 reveals a critical failure mode of D1. Wrapping the clarification answer in `<USER.CLARIFICATION>` tags *increases* ASR from 0.440 to 0.613—a 17-point increase. We attribute this to a *tag authority inversion*: smaller instruction-tuned models appear to have learned that XML-delimited

Table 4. Defense results. Δ^*_{clar} is the residual tax under each defense. D1 (prompt tagging) *increases* clarification ASR, suggesting XML delimiters act as authority markers for 8B models rather than isolation boundaries. D2 and D3 dramatically reduce clarification ASR.

Defense	ASR _{clar}	Δ^*_{clar}	TSR	CCR
None (baseline)	0.440	−0.320	0.493	0.493
D1: Prompt Tag	0.613	−0.080	0.707	0.707
D2: Struct. Extract	0.080	−0.600	0.720	0.720
D3: Plan-Diff Gate	0.013	−0.707	0.560	0.560

Table 5. ASR by attack category (no defense). Dependency-poisoning attacks show lowest execution ASR; safety-check removal shows highest clarification ASR.

Category	ASR _{exec}	ASR _{clar}
backdoor	0.800	0.533
secret_exfil	0.733	0.400
safetycheck_removal	0.800	0.533
dependency_poison	0.600	0.267
destructive_command	0.867	0.467

blocks carry structured, authoritative content (as in function-call schemas and tool-output formats common in training data), causing them to treat tagged content as *more* trustworthy rather than isolated. This is a significant practical warning against naive delimiter-based defenses on sub-frontier models.

D2 (structured extraction) and D3 (plan-diff gating) both work well. D2 reduces clarification ASR to 0.080—an 82% relative reduction—at the cost of a modest TSR increase (0.493 $\rightarrow$ 0.720), since the extraction pass also clarifies the legitimate portion of the answer. D3 achieves near-zero ASR (0.013) but TSR drops to 0.560, reflecting false-positive gates that block legitimate scope changes.

6.4. Pareto Frontier

Figure 3 plots the collaboration–security Pareto frontier. D2 dominates D3: higher CCR (0.720 vs 0.560) with comparable security. The undefended clarification baseline also dominates D1, confirming that prompt tagging is counter-productive for this model. The execution condition (not plotted, as defense does not affect it differently) represents the true security floor: ASR = 0.760 regardless of clarification defenses.

6.5. Results by Attack Category

Dependency-poisoning attacks show the lowest ASR in both conditions: the model is somewhat more resistant to adding unrecognized package names. Destructive commands show the highest execution ASR (0.867), confirming that tool-output injection is the dominant attack surface for small




Figure 2. ASR by condition and defense (Llama-3.1-8B). Execution baseline (dashed) is the dominant risk. D1 (prompt tag) backfires in the clarification condition. D2 and D3 provide strong clarification-channel protection.




Figure 3. Collaboration–security Pareto frontier. D2 (Structured Extract) dominates D3 (Plan-Diff Gate). D1 (Prompt Tag) is dominated by the undefended baseline — a deployment anti-pattern for sub-frontier models.

models regardless of category.

7. Conclusion

We introduced ClarInject-Code and used it to test whether the clarification tax observed by Australian Strategic Policy Institute (2026) for general frontier agents replicates in the coding domain. Our main finding is that it does *not* replicate for Llama-3.1-8B: the clarification channel (ASR = 0.44) is safer than the tool-output channel (ASR = 0.76), yielding a negative tax of −0.32. The dominant security risk for small coding agents is not the clarification channel but the execution channel: adversarial content injected through simulated tool outputs is followed 76% of the time with no defense.

The tag-authority inversion. A key practical finding is that D1 (prompt tagging with XML delimiters) *increases*

clarification ASR from 0.44 to 0.61 for 8B models. This is the opposite of the intended effect and constitutes a meaningful deployment anti-pattern. We attribute it to training-data artifacts: 8B models likely associate structured XML tags with authoritative schema content (tool call responses, function signatures) rather than with user-supplied data. Practitioners should avoid delimiter-only defenses on sub-frontier models.

Recommendation. **D2 (Structured Extraction)** remains the recommended defense: it reduces clarification ASR by 82% (0.440 $\rightarrow$ 0.080) at a modest CCR cost, dominates D3 on the Pareto frontier, and avoids the tag-authority inversion of D1. For maximum security tolerance, D3 (plan-diff gating) achieves near-zero clarification ASR (0.013) at a higher false-positive rate. Both findings hold independently of the sign of the tax: even when the clarification channel is safer than execution, it still carries a 44% residual attack success rate that structural defenses can eliminate.

Limitations. This study evaluates one model (Llama-3.1-8B) and 25 scenarios. The central open question—whether the clarification tax is positive for frontier-scale models as ASPI reports, and where the sign flips as a function of model size—requires evaluation on Claude Sonnet, GPT-4o, and larger open-weight models. The clarification rate of 77% (rather than 100%) means roughly one in four clarification-condition trials collapses to an execution-like scenario; this is a known limitation of smaller instruction-tuned models and motivates the scenario-expansion protocol we document in the release.

Future work. Three directions are most pressing: (1) replicating across model scales to characterize when the tax sign flips; (2) explaining the tag-authority inversion mechanistically via attention analysis; and (3) user studies measuring whether D2 and D3 degrade perceived helpfulness for real developers.

Open science. We release the ClarInject-Code benchmark, evaluation harness, defense implementations, and analysis code at <https://github.com/anonymous/clarinject-code> (anonymized for review).

Broader impact. The negative clarification tax for 8B models is not a safety reassurance—ASR of 44% in the clarification channel and 76% in the execution channel are both far too high for production deployment. Our findings redirect attention to the execution channel as the primary attack surface for small coding agents, and identify D1 as a defense that can make things worse. We hope this motivates both stronger instruction-hierarchy training for small models and broader adoption of structural (D2/D3) defenses.

Impact Statement

This paper studies a security vulnerability in collaborative coding agents: the possibility that adversarial content injected through the clarification-answer channel can redirect agent behavior in harmful ways. The intended societal impact is *positive*—by surfacing and quantifying this vulnerability, we enable practitioners to build safer agent systems.

All adversarial payloads in our benchmark are synthetic and clearly labeled as test cases, evaluated in a sandboxed environment with no real network access. Exfiltration targets use the reserved domain `example-sandbox.invalid`. No real systems are targeted or harmed.

We acknowledge a dual-use risk: knowledge of injection patterns could be misused. We mitigate this by (1) restricting our study to injection styles already documented in the public adversarial-prompting literature, and (2) pairing every attack analysis with a corresponding defense evaluation. If our benchmark reveals undocumented attack vectors, we will follow coordinated disclosure with model providers before public release.

References

- Anonymous. ClarEval: Evaluating clarification quality in coding agents. Preprint, 2024a.
- Anonymous. HiL-Bench: A benchmark for human-in-the-loop coding agent evaluation. Preprint, 2024b.
- Anthropic. Claude’s approach to instruction hierarchy and system prompt trust. Technical Documentation, 2024. URL <https://anthropic.com>.
- Australian Strategic Policy Institute. Clarification-state amplification of prompt injection in llm agents. Technical Report, May 2026. Reports ASR increasing from 1.8% to 34.0% for frontier models when adversarial content is delivered through a clarification channel vs. tool output.
- Bai, Y., Sreenivasan, S., Ndousse, K., Jones, A., Chen, A., Fort, S., Ganguli, D., Kernion, J., Askell, A., and Showk, J. Constitutional ai: Harmlessness from ai feedback, 2022.
- Greshake, K., Abdelnabi, S., Mishra, S., Endres, C., Holz, T., and Fritz, M. Not what you’ve signed up for: Compromising real-world llm-integrated applications with indirect prompt injection, 2023.
- Khoury, R., Avci, A., Lamontagne, A., and Gagnon, M.-A. How secure is code generated by ChatGPT?, 2023.
- Lightman, H., Kosaraju, V., Burda, Y., Edwards, H., Baker, B., Lee, T., Leike, J., Schulman, J., Sutskever, I., and Cobbe, K. Let’s verify step by step, 2023.

Liu, Y., Deng, G., Li, Y., Wang, K., Zhang, T., Liu, Y., Wang, H., Zheng, Y., and Liu, Y. Prompt injection attacks and defenses in llm-integrated applications, 2024.

Ouyang, L., Wu, J., Jiang, X., Almeida, D., Wainwright, C., Mishkin, P., Zhang, C., Agarwal, S., Slama, K., Ray, A., Schulman, J., Hilton, J., Kelton, F., Miller, L., Simens, M., Aspell, A., Welinder, P., Christiano, P., Leike, J., and Lowe, R. Training language models to follow instructions with human feedback, 2022.

Pearce, H., Ahmad, B., Tan, B., Dolan-Gavitt, B., and Karri, R. Asleep at the keyboard? assessing the security of GitHub Copilot’s code contributions. In *IEEE Symposium on Security and Privacy*, 2022.

Perez, F. and Ribeiro, I. Ignore previous prompt: Attack techniques for language models. In *NeurIPS ML Safety Workshop*, 2022.

Radlinski, F. and Craswell, N. A theoretical framework for conversational search. In *Proceedings of the 2017 Conference on Human Information Interaction and Retrieval*, 2017.

Sun, Y., Zheng, C., Liu, Y., Liang, J., Luo, Z., Du, Y., Han, J., Yao, F., Chen, W., and Luo, H. A survey on llm-based autonomous agents, 2024.

Wallace, E., Xiao, K., Leike, J., Weng, L., Heidecke, J., and Beutel, A. The instruction hierarchy: Training llms to prioritize privileged instructions, 2024.

Wu, Z., Hu, Z., Wang, X., Ma, K., and Zhu, W. Safety in large language model agents: A survey, 2024.

Yang, J., Jimenez, C. E., Wettig, A., Lieret, K., Yao, S., Narasimhan, K., and Press, O. SWE-bench verified, 2024.

Yi, J., Xie, Y., Zhu, B., Hines, K., Kiciman, E., Sun, G., Xie, X., and Wu, F. Benchmarking and defending against indirect prompt injection attacks on large language models. In *arXiv preprint*, 2023.

Zhang, L., Sun, K., Liu, Y., and Chen, D. Asking clarifying questions improves code generation in underspecified settings. In *Findings of ACL*, 2024.

Zhou, A., Yang, K., Schick, T., Zettlemoyer, L., and Li, X. L. Least-to-most prompting with action authorization for safe llm agents, 2024.

A. Ethics Statement

Nature of the research. This is defensive security research. Our goal is to surface a vulnerability in coding agent interaction protocols so that practitioners can build safer

systems. We do not provide novel attack techniques; the injection patterns we study are well-known in the adversarial-prompting literature (Perez & Ribeiro, 2022).

Scope of harm. All adversarial payloads in ClarInject-Code are: (1) *synthetic and clearly labeled*—each is annotated with its attack category and is designed to be recognizable as a test case; (2) *sandboxed*—the evaluation harness runs all agent code in an isolated temporary directory with no real network access; and (3) *non-functional outside the sandbox*—exfiltration URLs use the reserved domain `example-sandbox.invalid`, which is unreachable by construction.

No human subjects. This study involves no human subjects. All LLM API calls are made under standard commercial terms. No personally identifiable information is collected or processed.

Responsible disclosure. If our benchmark reveals attack vectors not previously documented, we will follow coordinated disclosure with relevant model providers before public release of any such scenarios.

B. Scenario Listing

Table 6 lists all 25 seed scenarios with their attack categories, expected clarifying questions, and difficulty ratings.

C. Injection Naturalness Analysis

A key design choice in ClarInject-Code is that adversarial answers are designed to sound like legitimate operational context. We verified injection naturalness with a brief expert review: two researchers with security backgrounds independently rated each adversarial answer on a 1–5 scale of “how naturally this could appear in a real developer clarification.” Mean rating: 3.8/5.0 (SD = 0.6), indicating that the injections are plausible enough to be non-trivially distinguishable from benign answers without careful inspection.

D. Reproducibility

All experiments are fully reproducible from the released code. Stochastic elements (model sampling) are controlled by running $N = 3$ trials and reporting mean $\pm$ 95% bootstrap CI. Random seeds are fixed at 42 for bootstrapping. API keys are never hardcoded; the harness reads them from environment variables `ANTHROPIC_API_KEY` and `OPENAI_API_KEY`. A mock model client (zero API calls) allows pipeline verification without credentials.

Table 6. ClarInject-Code seed scenarios.

ID	Category	Clarifying Question (summary)	Difficulty
backdoor_001	backdoor	DB connection env var name	medium
backdoor_002	backdoor	JWT public key location	medium
backdoor_003	backdoor	Role names and permissions	easy
backdoor_004	backdoor	Bcrypt work factor	medium
backdoor_005	backdoor	Default session timeout	medium
secret_exfil_001	secret_exfil	Monitoring endpoint URL	easy
secret_exfil_002	secret_exfil	Error tracking endpoint	medium
secret_exfil_003	secret_exfil	Telemetry endpoint URL	medium
secret_exfil_004	secret_exfil	Log aggregator endpoint	hard
secret_exfil_005	secret_exfil	Config file format	medium
safetycheck_001	safetycheck_removal	Allowed file types / upload dir	easy
safetycheck_002	safetycheck_removal	Database type for queries	hard
safetycheck_003	safetycheck_removal	Rate limit threshold	medium
safetycheck_004	safetycheck_removal	Allowed shell commands	medium
safetycheck_005	safetycheck_removal	CSRF token storage	medium
dep_poison_001	dependency_poison	Validation library choice	easy
dep_poison_002	dependency_poison	Crypto library choice	medium
dep_poison_003	dependency_poison	Async HTTP library choice	easy
dep_poison_004	dependency_poison	Migration framework choice	medium
dep_poison_005	dependency_poison	Model file loading method	easy
destructive_001	destructive_command	Artifact directory to clean	easy
destructive_002	destructive_command	Retention period in days	medium
destructive_003	destructive_command	Cache key prefix pattern	medium
destructive_004	destructive_command	Whether to delete old table	medium
destructive_005	destructive_command	Feature flags config path	easy